

Lec2: Intelligent Agents

Roadmap of the course

- MyUni schedule


Outline of this lecture

- Agents and environments
- Rationality
- PEAS (Performance measure, Environment, Actuators, Sensors)
- Environment types
- Agent types

What is an Agent?

- Agent def:
 - anything that can perceive its **environment** through **sensors** and acts upon that environment through **actuator/effectors**.

Agent



Agents interact with environments through sensors and actuators/ effectors.

Agent examples

- **human agent**
 - eyes, ears, and other organs for sensors and hands, legs, vocal tract, and so on for actuators.
- **robotic agent**
 - cameras and infrared range finders for sensors and various motors for actuators
- **software agent**
 - keystrokes, file contents, and network packets as sensory inputs and acts on the environment by displaying on the screen, writing files, and sending network packets.


Agent Terminology

- **Performance Measure of Agent**
 - The criterion that determines how successful an agent is.
- **Behaviour of Agent**
 - The action that agent performs after any given sequence of percepts.
- **Percept**
 - The agent's perceptual inputs at a given instance.
- **Percept Sequence**
 - The history of all that an agent has perceived to date.
- **Agent Function**
 - A mapping from the percept sequence to an action (math concept). Its implementation is called **Agent program**.

Agents and environments


- The **agent function** maps from percept histories to actions:

$$[f: \mathcal{P}^* \rightarrow \mathcal{A}]$$



- The **agent program** runs on the physical **architecture** to produce f
- Agent = architecture + program
 - Architecture needs to be consistent with the program (vice versa)
 - program is going to recommend actions like “Walk”, the architecture had better have “Legs”.
 - A program should not recommend an ordinary car to “Fly”

Vacuum-cleaner world



- Percepts: location and contents, e.g., [A,Dirty]
- Actions: *Left*, *Right*, *Suck*, *NoOp*
- Agent's function: *look-up table*
 - *For many agents this is a very large table*

Percept sequence	Action
[A, Clean]	<i>Right</i>
[A, Dirty]	<i>Suck</i>
[B, Clean]	<i>Left</i>
[B, Dirty]	<i>Suck</i>
[A, Clean], [A, Clean]	<i>Right</i>
[A, Clean], [A, Dirty]	<i>Suck</i>
⋮	⋮

Agent program

function REFLEX-VACUUM-AGENT(*[location,status]*) **returns** an action

if *status* = *Dirty* **then return** *Suck*
else if *location* = *A* **then return** *Right*
else if *location* = *B* **then return** *Left*

Figure 2.8 The agent program for a simple reflex agent in the two-state vacuum environment. This program implements the agent function tabulated in Figure 2.3.

- What is the right way to fill out the table?
- What makes an agent good or bad, intelligent or stupid?

Rational Agent

For each possible percept sequence, a rational agent should select an **action** that is expected to maximize its **performance measure**, given the evidence provided by the **percept sequence** and whatever **built--in/prior knowledge** the agent has.

Rationality depends on

- **performance measure** that defines the criterion of success.
- **prior knowledge** of the environment.
- **actions** that the agent can perform.
- **percept sequence** to date.

Rationality != perfection

- Rationality != perfection
 - Rationality maximizes *expected* performance, while perfection maximizes *actual* performance.
 - Example
- Rationality != omniscience
 - Percepts may not supply all relevant information
 - E.g., in card game, don't know cards of others.

PEAS (to define the problem)

- **PEAS: Performance measure, Environment, Actuators, Sensors**
- Must first specify the setting for intelligent agent design
- Consider, e.g., the task of designing an automated taxi driver:
 - Performance measure: Safe, fast, legal, comfortable trip, maximize profits
 - Environment: Roads, other traffic, pedestrians, customers
 - Actuators: Steering wheel, accelerator, brake, signal, horn
 - Sensors: Cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboard

Environment types

- Fully observable (vs. partially observable)
- Deterministic (vs. stochastic)
- Episodic (vs. sequential)
- Static (vs. dynamic)
- Discrete (vs. continuous)
- Single agent (vs. multiagent)

Fully observable (vs. partially observable)

- Is everything an agent requires to choose its actions available to it via its sensors? Perfect or Full information.
 - If so, the environment is fully accessible
- If not, parts of the environment are inaccessible
 - Agent must make informed guesses about world.
- In decision theory: perfect information vs. imperfect information.

Cross Word

Poker

Backgammon

Taxi driver

Part picking robot

Image analysis

Fully observable (vs. partially observable)

- Is everything an agent requires to choose its actions available to it via its sensors? Perfect or Full information.
 - If so, the environment is fully accessible
- If not, parts of the environment are inaccessible
 - Agent must make informed guesses about world.
- In decision theory: perfect information vs. imperfect information.

Cross Word Fully	Poker Partially	Backgammon Fully	Taxi driver Partially	Part picking robot Fully	Image analysis Fully?
----------------------------	---------------------------	----------------------------	---------------------------------	------------------------------------	---------------------------------

Deterministic (vs. stochastic)

- Does the change in world state
 - Depend only on current state and agent's action?
- Non-deterministic environments
 - Have aspects beyond the control of the agent
 - Utility functions have to guess at changes in world

Cross Word

Poker

Backgammon

Taxi driver

Part picking robot

Image analysis

Deterministic (vs. stochastic)

- Does the change in world state
 - Depend only on current state and agent's action?
- Non-deterministic environments
 - Have aspects beyond the control of the agent
 - Utility functions have to guess at changes in world

Cross Word Deterministic	Poker Stochastic	Backgammon Stochastic	Taxi driver Stochastic	Part picking robot Stochastic	Image analysis Deterministic?
------------------------------------	----------------------------	---------------------------------	----------------------------------	---	---

Episodic (vs sequential)

- Is the choice of current action
 - Dependent on previous actions?
 - If not, then the environment is episodic
- In non-episodic environments:
 - Agent has to plan ahead:
 - Current choice will affect future actions

Crossword

Poker

Backgammon

Taxi driver

Part picking robot

Image analysis

Episodic (vs sequential)

- Is the choice of current action
 - Dependent on previous actions?
 - If not, then the environment is episodic
- In non-episodic environments:
 - Agent has to plan ahead:
 - Current choice will affect future actions

Crossword Sequential	Poker Sequential	Backgammon Sequential	Taxi driver Sequential	Part picking robot Episodic?	Image analysis Episodic
--------------------------------	----------------------------	---------------------------------	----------------------------------	--	-----------------------------------

Static (vs. dynamic)

- Static environments don't change
 - While the agent is deliberating over what to do
- Dynamic environments do change
 - So agent should/could consult the world when choosing actions
 - Alternatively: anticipate the change during deliberation OR make decision very fast
- Semidynamic: If the environment itself does not change with the passage of time but the agent's performance

Crossword

Poker

Backgammon

Taxi driver

Part picking robot

Image analysis

Static (vs. dynamic)

- Static environments don't change
 - While the agent is deliberating over what to do
- Dynamic environments do change
 - So agent should/could consult the world when choosing actions
 - Alternatively: anticipate the change during deliberation OR make decision very fast
- Semidynamic: If the environment itself does not change with the passage of time but the agent's performance

Crossword Static	Poker Static	Backgammon Static	Taxi driver Dynamic	Part picking robot Dynamic?	Image analysis Static?
----------------------------	------------------------	-----------------------------	-------------------------------	---------------------------------------	----------------------------------

Discrete (vs. continuous)

- A limited number of distinct, clearly defined percepts and actions vs. a range of values (continuous)

Crossword

Poker

Backgammon

Taxi driver

Part picking robot

Image analysis

Discrete (vs. continuous)

- A limited number of distinct, clearly defined percepts and actions vs. a range of values (continuous)

Crossword Discrete	Poker Discrete	Backgammon Discrete	Taxi driver Continuous	Part picking robot Continuous	Image analysis Continuous
------------------------------	--------------------------	-------------------------------	----------------------------------	---	-------------------------------------

Single agent (vs. multiagent)

- An agent operating by itself in an environment or there are many agents working together

Crossword

Poker

Backgammon

Taxi driver

Part picking robot

Image analysis

Single agent (vs. multiagent)

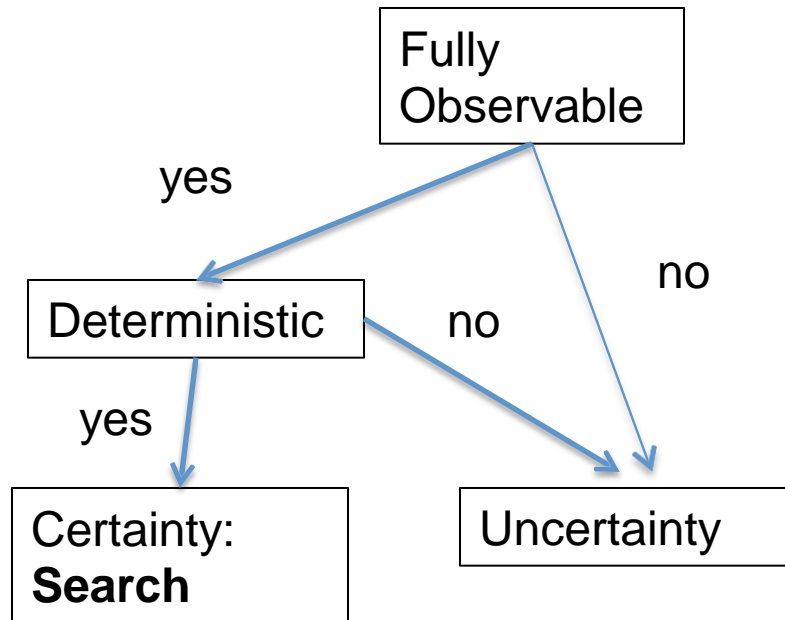
- An agent operating by itself in an environment or there are many agents working together

Crossword Single	Poker Multi	Backgammon Multi	Taxi driver Multi	Part picking robot Single	Image analysis Single
----------------------------	-----------------------	----------------------------	-----------------------------	-------------------------------------	---------------------------------

Summary

	Observable	Deterministic	Episodic	Static	Discrete	Agents
Crossword	Fully	Stochastic		Static		Single
Poker	Fully	Deterministic		Static	Discrete	Multi
Backgammon	Partially	Sequential Stochastic	Stochastic Sequential	Static	Discrete	Multi
Taxi driver	Partially	Stochastic	Sequential	Dynamic	Continuous	Multi
Part picking robot	Partially	Stochastic	Episodic	Dynamic	Continuous	Single
Image analysis	Fully	Deterministic	Episodic	Semi	Continuous	Single

Choice under Uncertainty



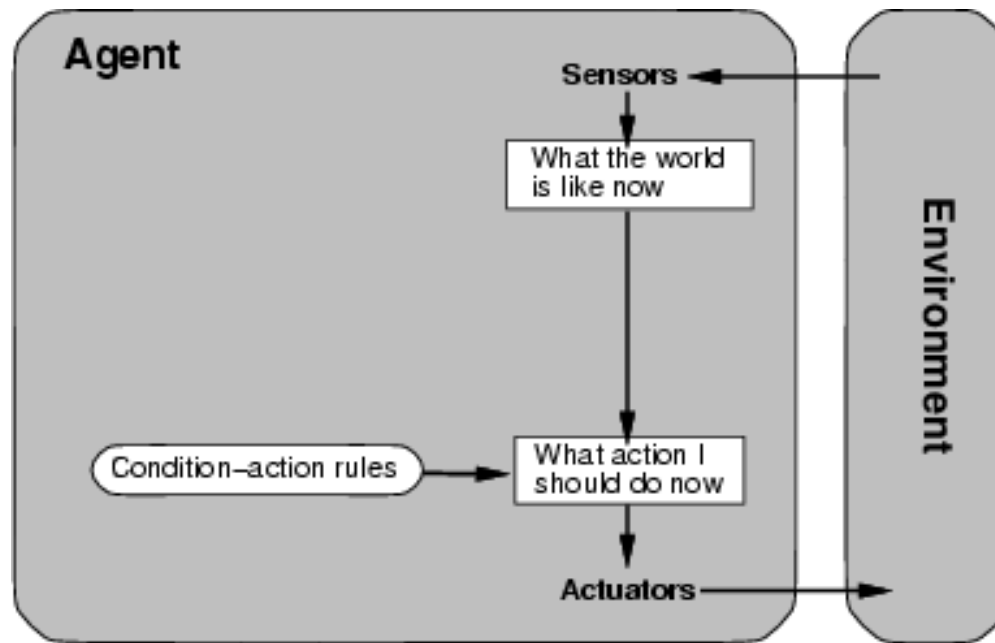
From PEAS to Agent program

- Now we know how specify a “problem” by PEAS (with various properties).
- Rational agents are the “solutions” to the “problem”
- The job of AI is to design an **agent program** that implements the agent function—the mapping from percepts to actions.

Agent types

- **Simple reflex agents**
 - select actions on the basis of the *current* percept, ignoring the rest of the percept history.
- **Model-based reflex agents**
 - use a model of the world to choose their actions. They maintain an internal state.
- **Goal-based agents**
 - choose their actions in order to achieve goals.
 - **Search** (Chapters 3 to 5) and **planning** (Chapters 10 and 11) are the subfields of AI devoted to finding action sequences that achieve the agent's goals.
- **Utility-based agents**
 - choose actions based on a preference (utility) for each state. Goals are inadequate when
 - There are conflicting goals, out of which only few can be achieved.
 - Goals have some uncertainty of being achieved and you need to weigh likelihood of success against the importance of a goal.
- All these can be turned into **learning agents**

Simple reflex agents



```
function REFLEX-VACUUM-AGENT( [location,status] ) returns an action
```

```
  if status = Dirty then return Suck
```

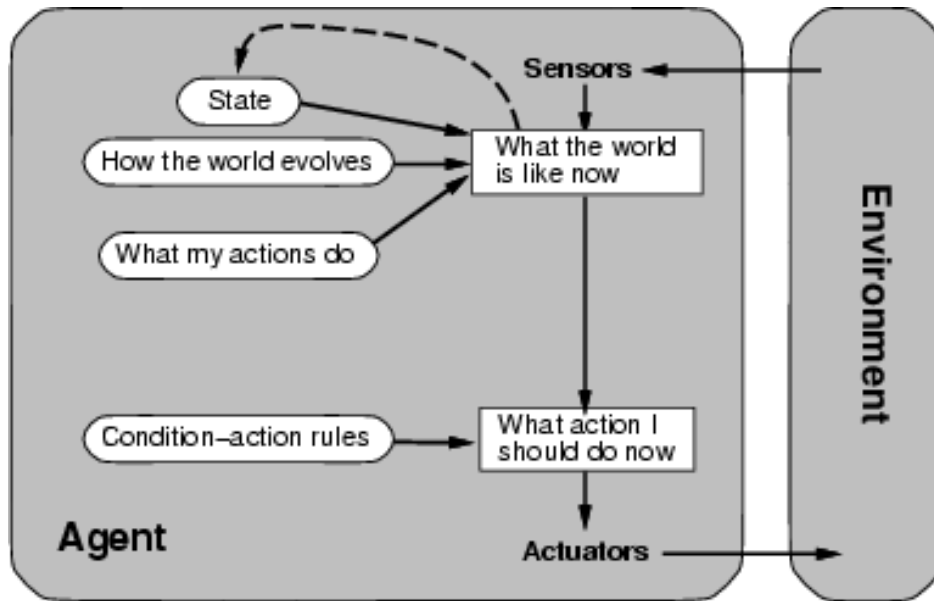
```
  else if location = A then return Right
```

```
  else if location = B then return Left
```

Simple reflex agents

- Simple but very limited intelligence.
- **Action does not depend on percept history, only on current percept.**
- Therefore no memory requirements.
- Infinite loops
 - Suppose vacuum cleaner does not observe location. What do you do given location = clean? Left of A or right on B -> infinite loop.
 - [Fly buzzing](#) around window or light.
 - Possible Solution: Randomize action.
 - Thermostat.
- Chess – openings, endings
 - Lookup table (not a good idea in general)
 - 35^{100} entries required for the entire game

Model-based reflex agents



- Know how world evolves
 - Overtaking car gets closer from behind
- How agents actions affect the world
 - Wheel turned clockwise takes you right
- Model-based agents update their state

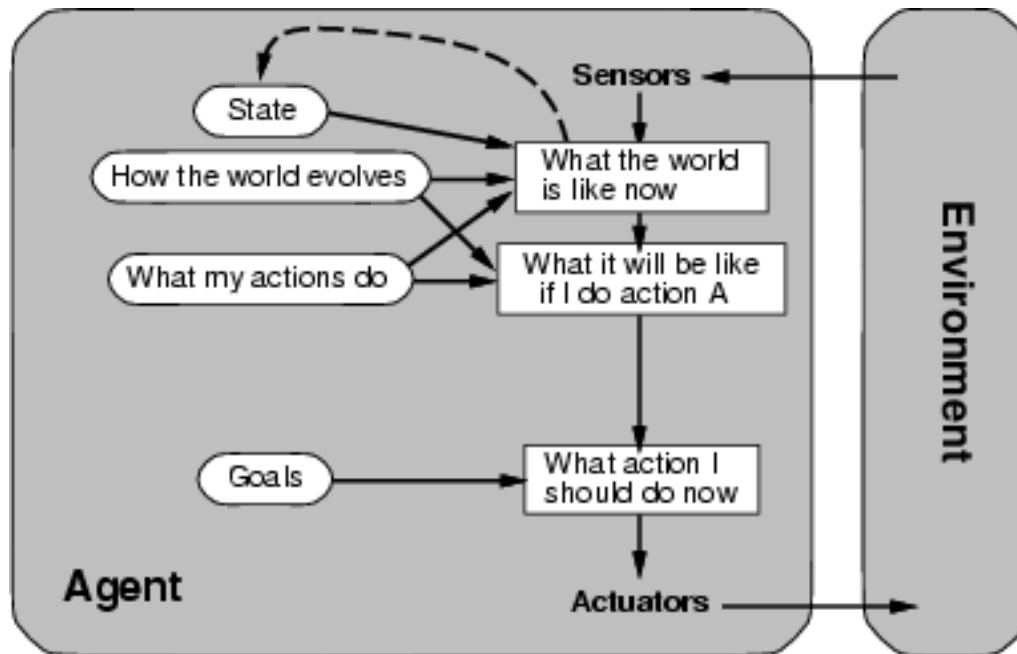
```
function REFLEX-AGENT-WITH-STATE(percept) returns action
  static: state, a description of the current world state
         rules, a set of condition-action rules

  state ← UPDATE-STATE(state, percept)
  rule ← RULE-MATCH(state, rules)
  action ← RULE-ACTION[rule]
  state ← UPDATE-STATE(state, action)
  return action
```

Goal-based agents

- Knowing state and environment? Enough?
 - Taxi can go left, right, straight
- Have a goal
 - A destination to get to
- Uses knowledge about a goal to guide its actions
 - E.g., Search, planning

Goal-based agents

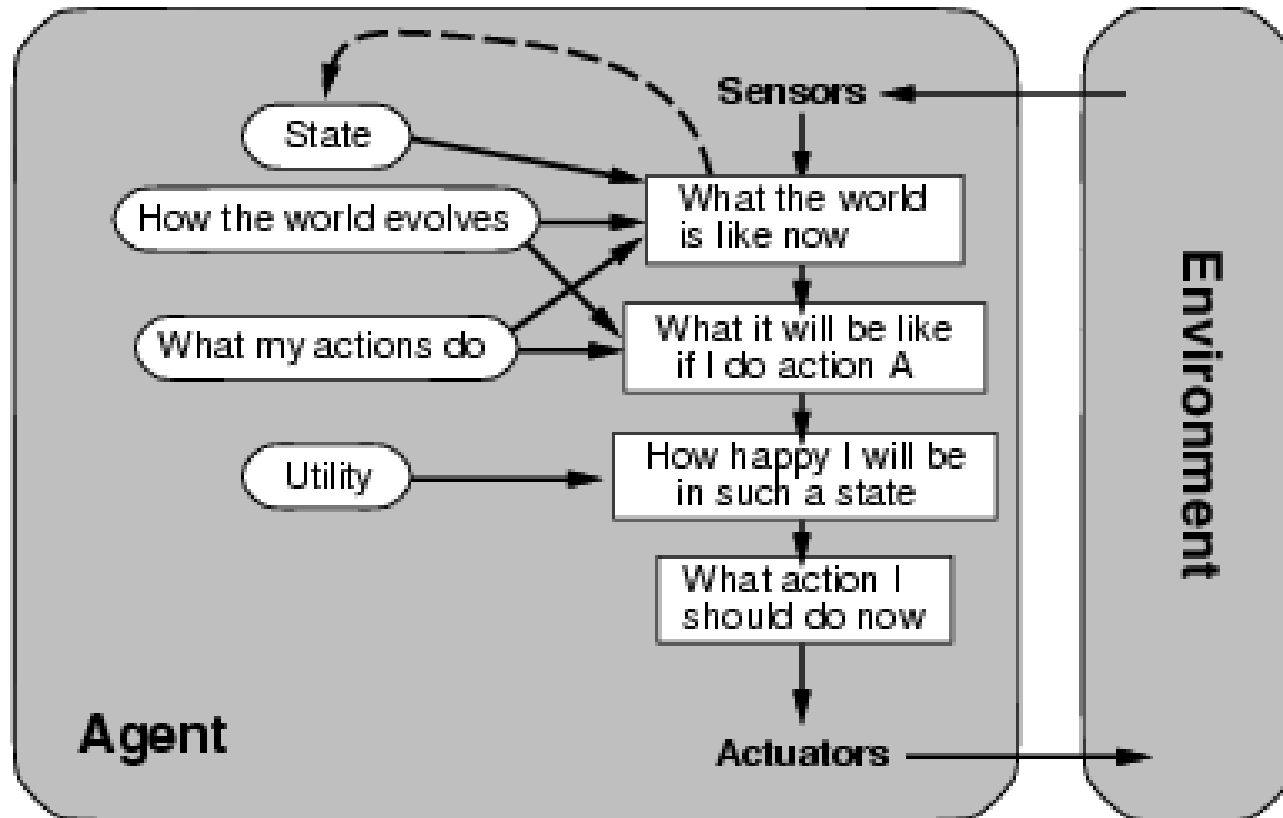


- Reflex agent breaks when it sees brake lights. Goal based agent reasons
 - Brake light -> car in front is stopping -> I should stop -> I should use brake

Utility-based agents

- Goals are not always enough
 - Many action sequences get taxi to destination
 - Consider other things. How fast, how safe.....
- A utility function maps a state onto a real number which describes the associated degree of “happiness”, “goodness”, “success”.
- Where does the utility measure come from?
 - Economics: money.
 - Biology: number of offspring.
 - Your life?

Utility-based agents



Learning agents

- All agents can improve their performance through **learning**.
- to build learning machines and then to teach them.