

Mining Big Data

Finding Similar Items (Chapter 3)

Motivation

- Finding similar items in a set of documents is a fundamental problem in data mining.

Example:

- Given a collection of web pages, the goal is to find near-duplicate pages.
- Such pages could be plagiarisms or mirrors

An important problem that arises is that there may be far too many pairs of items to test for similarity


Applications of Near-Neighbor Search

- We need a notion of similarity.
- We measure the similarity of sets by the relative size of their intersection
- This is called “Jaccard similarity”.
- Applicable to
 - finding textually similar documents
 - collaborative filtering by finding similar customers and similar products

Jaccard Similarity of Sets

- Given two sets S and T , the Jaccard similarity is defined as $|S \cap T| / |S \cup T|$ (size of the intersection divided by size of the union)

Example with Jaccard similarity 3/8.
Figure 3.1 in Rajaramam/Ullman



Similarity of Documents

- Jaccard similarity works well for finding textually similar documents in a large set of documents (for example the Web or collection of news articles)
- We are currently looking at “character level” similarity not “similar meaning”
- Testing whether two documents are the same is easy (compare documents by character).

Examples

For many documents large portions of the text are identical.

Examples:

- Plagiarism
- Mirror Pages
- Articles from the same source

Collaborative Filtering

- Collaborative filtering is another application of similarity of sets.
- Goal is to recommend to users items that were liked by other users who have exhibited similar tastes.

Examples:

- Online Purchases
- Movie Ratings

Example: Online Purchases

- Amazon has millions of customers and sells millions of items
- Its database records which items have been bought by which customers.
- We can say that two customers are similar if their sets of purchased items have high Jaccard similarity
- Likewise, two items are similar if the sets of purchasers have high Jaccard similarity.
- We might expect mirror sites to have a Jaccard similarity of 90%
- For customers with similar tastes this will be much lower (20% might be unusual).
- Collaborative filtering is often combined with clustering (see Chapter 7) in this case.

Movie Ratings

- Netflix records which movies each of its customers rented and ratings of movies by customers.
- We can regard movies as similar if rented or rated highly by many of the same customers.
- We can regard customers as similar if they rented or rated highly many of the same movies.
- Similarities don't need to be high to be significant (same as for Amazon)
- Clustering movies by genre will make things easier.

Movie Ratings

How to deal with ratings.

- Ignore low-rated customer/movie pairs and treat them as if the customer never rented the movie.
- Two elements for each movie: liked or hated. Use Jaccard similarities for these different categories.
- If ratings are 1-to-5 stars: Put a movie in a customers set i times if they rated it i stars. Use Jaccard similarity for bags B and C and count an element j times in the intersection if j is the minimum of the number of times the element appears in B and C .

Shingling of Documents

- Represent documents as sets by constructing the set of short strings that appear in a document.
- Documents that share pieces (sentences, phrases) will have many common elements in their sets.

K-Shingles

- A document is a string of characters.
- A k-shingle for a document is any substring of length k in that document
- We can associate with each document its set of k-shingles (appear at least once in the document).

Example:

- Document D is the string abcdabd
- Set of 2-shingles for D is {ab,bc,cd,da,bd}
- Note that ab appear twice in D (but not in 2-shingles)

White Spaces

Several ways to treat white spaces (blank, tab, newline, etc).

Makes sense to replace any sequence of white spaces by a single blank.

Example:

- If we use $k=9$ and eliminate all white spaces then “The plane was ready for **touch down**” and
- “The quarterback scored at **touchdown**” look similar.
- Doesn’t hold if we keep a blank.

Shingle Size

- We can pick k as any constant we like.
- If k is too small, sequences of k characters appear in most documents.
- We could have documents whose shingle-sets have high Jaccard similarity although the documents don't share phrases or sentences.
- Good choice of k depends on how long typical documents are and how large the set of typical characters is.
- k should be large enough such that the probability of any given shingle in any given document is low.

Shingle Size

- If corpus of documents is emails, $k=5$ should be fine.
- Suppose that only letters and general white-spaces appear.
- This implies $27^5=14,348,907$ possible shingles.
- Typical email is much smaller than 14 million characters and we expect $k=5$ to work well (and it does).
- Calculation is more subtle as characters don't appear with equal probability (for example letter "z" doesn't appear to often)
- Good rule of thumb is to imagine that there are only 20 characters and estimate number of k -shingles as 20^k .
- For large documents, such as research articles, $k=9$ is considered safe.

Hashing Shingles

- Instead of using substrings directly, we can pick a hash function and map a string of length k to some number of buckets
- Treat the resulting bucket as shingle.
- Set representing a document is the set of integers that are bucket numbers of one or more k -shingles that appear in the document.

Example:

- Set of 9-shingles for a document
- Map each 9-shingle to a bucket number in the range 0 to $2^{32}-1$.
- Each shingle is represented by 4 bytes (instead of nine).

Data is compacted and we can manipulate (hashed) shingles by single-word machine operations.

Shingle from Words

- We want to identify similar news articles.
- News articles are most prose and have a lot of stop words such as “and”, “you”, “to”.
- Often we ignore stop words since they are not useful.
- For finding similar news articles, it has been shown that using shingles defined by a stop word plus the next two words is useful.

Example

- An ad might have the simple text “Buy Sudzo”
- A news article with the same idea might read as “A spokesperson **for the** Sudzo Corporation revealed today **that** studies **have** shown **it is good for** people **to** buy Sudzo products” (stop words in red)
- First three shingles made from stop word + 2 following words are:
 - A spokesperson for
 - For the Sudzo
 - the Sudzo Corporation
- In total nine such shingles from the sentence, but none from the ad.

Similarity-Preserving Summaries

- Sets of shingles are large (even holds if we hash them, 4 times the size of the document)
- If we have millions of documents, it might not be possible to store all the shingle-sets in main memory.
- We want to replace large sets by much smaller representations called “signatures”.
- Import property: We can compare the signatures of two sets and estimate the Jaccard similarity of the underlying sets.
- Signatures can’t give exact similarity.
- But they provide good estimates and get more accurate for larger signatures.

Matrix Representation of Sets

- Goal is to get good signatures, let's start with representations for sets.
- We now visualize a collection of sets by their characteristic matrix.
- Columns correspond to the sets, rows correspond to the elements.
- $(r,c)=1$ if element for row r is member of set for column c .
- Otherwise the value in position $(r,c)=0$.
- If rows are products and columns are customers, the matrix represents the products bought by the different customers.

Example

Matrix representing four sets.

<i>Element</i>	S_1	S_2	S_3	S_4
<i>a</i>	1	0	0	1
<i>b</i>	0	0	1	0
<i>c</i>	0	1	0	1
<i>d</i>	1	0	1	1
<i>e</i>	0	0	1	0

Figure 3.2 in Rajaramam/Ullman

Minhashing

- Signatures that we want to construct for sets are composed of results of a large number of calculations (say several hundreds).
- Each calculation is a “minhash” of the characteristic matrix.
- We want to see how minhash is computed in principle.
- To minhash a set represented by a column of the characteristic matrix, pick a permutation of the rows.
- The minhash value of any column is the number of the first row (in permuted order) in which the column has a 1.

Example

- Suppose we pick order of rows beadc.
- This permutation defines a minhash function h .

- We have

– $h(S_1)=a$

– $h(S_2)=c$

– $h(S_3)=b$

– $h(S_4)=a$

<i>Element</i>	S_1	S_2	S_3	S_4
b	0	0	1	0
e	0	0	1	0
a	1	0	0	1
d	1	0	1	1
c	0	1	0	1

Figure 3.3 in Rajaramam/Ullman

MinHashing and Jaccard Similarity

Connection between minhashing and Jaccard similarity:

- The probability that the minhash function for a random permutation of the rows produces the same value for two sets equals the Jaccard similarity of those sets.

MinHashing and Jaccard Similarity

Reason: Consider the columns for those two sets, let's say S_1 and S_2 .

- Rows can be divided into three classes:
 - Type X rows have 1 in both columns
 - Type Y rows have 1 in one of the columns and 0 in the other
 - Type Z rows have 0 in both columns
- Consider similarity $SIM(S_1, S_2)$ and probability that $h(S_1) = h(S_2)$
- Let there be x rows of type X and y rows of type Y.
- We have $SIM(S_1, S_2) = x/(x+y)$.
- Consider rows that are permuted randomly.
- We have $h(S_1) = h(S_2)$ if we meet a X row before a Y row.
- When moving from top to bottom the **probability** of meeting a type X row before a type Y row is $x/(x+y)$.

Minhash Signatures

- Think of collection of sets by characteristic matrix M .
- To represent sets, we pick at random some number n of permutations of rows of M .
- 100 permutations or several hundred permutations will do.
- We call the minhash functions by these permutations h_1, h_2, \dots, h_n
- From the column representing set S , construct the minhash signature for S given by the vector $[h_1(S), h_2(S), \dots, h_n(s)]$.
- We can form from M a signature matrix in which the i th column is replaced by the minhash signature.
- Resulting matrix has same number of columns as M , but only n rows (much smaller than M).

Computing Minhash Signatures

- It's not feasible to permute a large characteristic matrix explicitly.
- It's possible to simulate the effect of a random permutation by a random hash function that maps row numbers to as many buckets as there are rows.
- Hash function that maps integers $0, 1, \dots, k-1$ to bucket number $0, 1, \dots, k-1$ will typically lead to collisions and leave other buckets unfilled.
- However, difference is unimportant as long as k is large and there are not too many collisions.

Computing Minhash Signatures

- Instead of picking n random permutations of rows, we pick n randomly chosen hash functions h_1, h_2, \dots, h_n on the rows.
- We construct signature matrix by considering each row r in given order
 1. Compute $h_1(r), h_2(r), \dots, h_n(r)$.
 2. For each column c do the following:
 - (a) If c has 0 in row r , do nothing.
 - (b) However, if c has 1 in row r , then for each $i = 1, 2, \dots, n$ set $SIG(i, c)$ to the smaller of the current value of $SIG(i, c)$ and $h_i(r)$.

Figure 3.4 in Rajaramam/Ullman

Example (3.8 in Rajaramam/Ullman)

- Consider the hash functions $h1(x) = x+1 \bmod 5$ and $h2(x) = 3x + 1 \bmod 5$.
- Letters naming the rows are replaced by numbers 0 through 4 and hash are applied to them.

<i>Row</i>	S_1	S_2	S_3	S_4	$x + 1 \bmod 5$	$3x + 1 \bmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

Example

Initial state: Matrix consist of all ∞ 's.

	S_1	S_2	S_3	S_4
h_1	∞	∞	∞	∞
h_2	∞	∞	∞	∞

Row 0: 1s in S_1 and S_4 and $h_1(0)=h_2(0)=1$.

This leads to

	S_1	S_2	S_3	S_4
h_1	1	∞	∞	1
h_2	1	∞	∞	1

Example

Row 1: 1 in S_3 and $h_1(1)=2$, $h_2(1)=4$.

- We set $SIG(1,3)=2$, $SIG(2,3)=4$

	S_1	S_2	S_3	S_4
h_1	1	∞	2	1
h_2	1	∞	4	1

Row2: 1s in S_2 and S_4 , $h_1(2)=3$, $h_2(2)=2$.

	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	1	2	4	1

Example

Row 3: 1s in S_1, S_3, S_4 . $h_1(3)=4$, $h_2(3)=0$.

	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	0	2	0	0

Row 4: 1 in S_3 , $h_1(4)=0$, $h_2(4)=3$.

	S_1	S_2	S_3	S_4
h_1	1	3	0	1
h_2	0	2	0	0

Example

- We can estimate the Jaccard similarities of the underlying sets from the signature matrix.

	S_1	S_2	S_3	S_4
h_1	1	3	0	1
h_2	0	2	0	0

We can estimate the Jaccard similarity by the fraction of rows where given two sets agree.

- $\text{SIM}(S_1, S_4) = 1.0$, but the true Jaccard similarity of S_1 and S_4 is $2/3$.
- $\text{SIM}(S_1, S_3) = 1/2$ (true $1/4$)
- $\text{SIM}(S_1, S_2) = 0$ (correct)

Locality-Sensitive Hashing

- Minhashing compresses large documents into small signatures and preserves expected similarity of any pair of documents.
- It still might be impossible to find the pairs with greatest similarity efficiently.

Reason:

- Number of pairs of documents may be too large even if there are not too many documents.

Example

- Assume we have a million documents and use signatures of length 250.
- We use 1000 bytes per document for signatures and entire data fits in a gigabyte (and therefore in main memory)
- However there are roughly $1,000,000^2/2$ (half a trillion) pairs of documents.
- If it takes a microsecond to compute the similarity of two signatures, then it takes almost six days to compute all similarities (on a laptop)

Local-sensitivity Hashing

- If we want to compute the similarity of every pairs, there is nothing we can do (except parallelizing the computation of similarities)
- Often we want the most similar pairs or all pairs above some lower bound in similarity.
- Local-sensitivity hashing (LSH) allows to do this without investigating every pair.

Procedure for Finding Similar Documents:

- Pick value of k and construct from each document the set of k -shingles.
- Sort document-shingle pairs by shingle
- Pick length n of minhash signatures and compute minhash signatures for all documents.
- Choose a threshold t that defines similarity.
- Pick number of bands b and number of rows r such that $br=n$ and threshold for S-curve is lower than t (limits false negatives).
- Construct candidate pairs by LSH
- Examine each candidate pair and determine whether similarity is at least t .
- **Optional:** If signatures are sufficient similar, check the documents directly.

LSH for Minhash Signatures

- One approach to LSH is to hash items several times.
- Similar items are more likely to be hashed to the same bucket than dissimilar ones.
- We consider any pair that hashed to the same bucket for any of the hashings as a candidate pair.
- Check only candidate pairs for similarity
- Dissimilar pairs that are hashed to the same bucket are false-positive (we hope that there are not too many)
- Truly similar pairs are missed (false-negative) if they are not hashed to the same bucket for a least one of the hash function (hope that there are just a few)

- If we have minhash signatures for the items, we can divide the signature matrix into b bands consisting of r rows each.
- For each band, there is a hash function that takes the vectors of r integers and hashes them to some large number of buckets.
- We can use the same hash function for each band, but have to use separate buckets for each band.

Bands

- Dividing signature matrix into bands

band 1	...	1 0 0 0 2	...
band 2		3 2 1 2 2	
band 3		0 1 3 1 1	
band 4			

Figure 3.6 in Rajaramam/Ullman

Analysis of Banding

- Suppose that we use b bands of r rows each and that a particular pair of documents has Jaccard similarity s .
- The probability that the minhash signatures for these documents agree in any particular row is s .

Probability that these documents become candidate pair:

- Probability that the signatures agree in all rows of one particular bands is s^r .
- Probability that they do not agree in at least one row is $1-s^r$.
- Probability that they do not agree in all rows of any of the bands is $(1-s^r)^b$.
- Probability that the signatures agree in all rows of at least one band (and mapped to same bucket) is $1-(1-s^r)^b$.

S-Curve




Figure 3.7 in Rajaramam/Ullman

- Regardless of chosen constants b and r , the probability function has the form of an S-curve. (steepest rise is determined by b and r)
- $(1/b)^{(1/r)}$ is approximation of threshold.
- Example $b=16$, $r=4$, threshold is approximately $\frac{1}{2}$.

Example

- Consider $b=20$, $r=5$ (signatures of length 100)
- Functions values for $1-(1-s^r)^b = 1-(1-s^5)^{20}$

s	$1 - (1 - s^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996

Figure 3.8 in Rajaramam/Ullman

Combining Techniques

- We can combine the examined techniques to find similar documents.
- The approach can produce false negatives (pairs of similar documents that are not identified).
- It also contains false positives (pairs that are evaluated but not found to be sufficiently similar)

Procedure for Finding Similar Documents:

- Pick value of k and construct from each document the set of k -shingles.
- Sort document-shingle pairs by shingle
- Pick length n of minhash signatures and compute minhash signatures for all documents.
- Choose a threshold t that defines similarity.
- Pick number of bands b and number of rows r such that $br=n$ and threshold for S-curve is lower than t (limits false negatives).
- Construct candidate pairs by LSH
- Examine each candidate pair and determine whether similarity is at least t .
- **Optional:** If signatures are sufficient similar, check the documents directly.